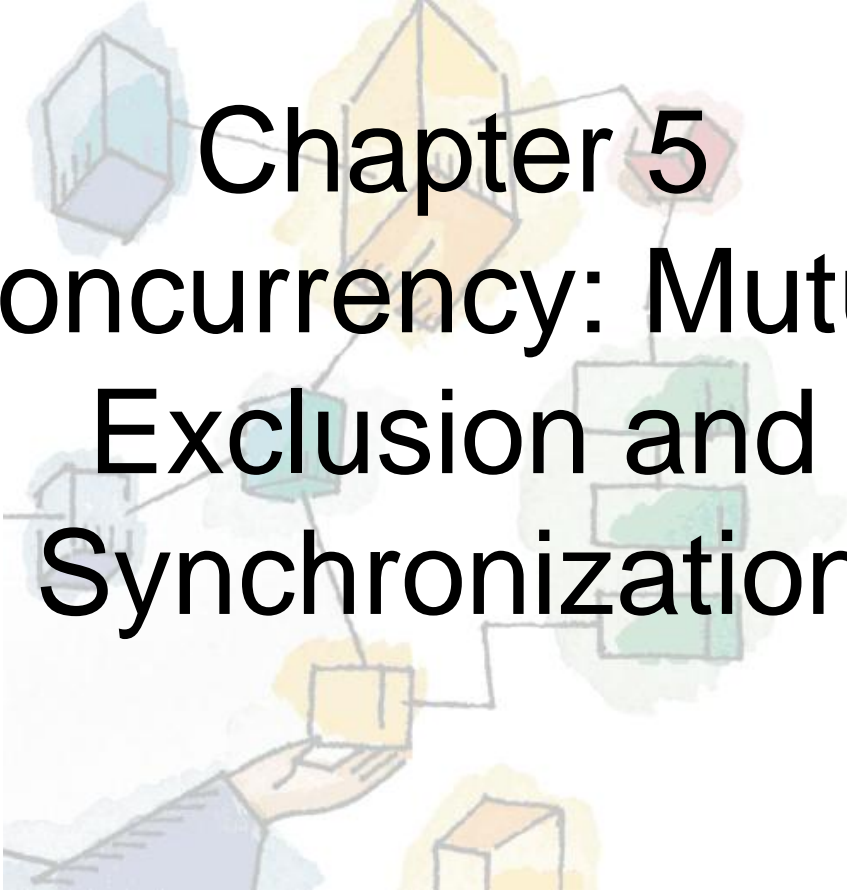


*Operating Systems:
Internals and Design Principles, 6/E*
William Stallings



Chapter 5

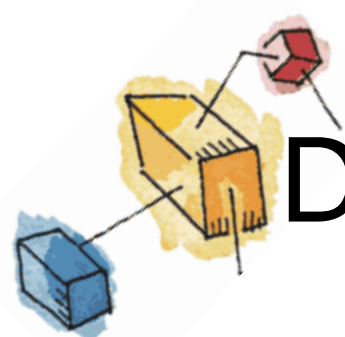
Concurrency: Mutual Exclusion and Synchronization



Multiple Processes

- Central to the design of modern Operating Systems is **managing multiple processes**
 - Multiprogramming
 - Multiprocessing
 - Distributed Processing
- Big Issue is **Concurrency**
 - Managing the interaction of all of these processes

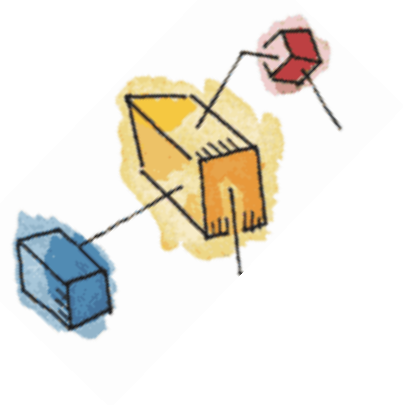




Difficulties of Concurrency

- **Sharing of global resources**
 - If P1 and P2 are sharing **common variable**, then care should be taken with read and write operations
- **Optimally managing the allocation of resources**
 - If P1 is given access to **I/O channel** and it gets suspended before using I/O, then performance will suffer
- **Difficult to locate programming errors as results are not deterministic and reproducible**
 - Due to overwriting of values





A Simple Example

```
void echo()
{
    chin = getchar();
    chout = chin;
    putchar(chout);
}
```

- Consider that we have a **single-processor multiprogramming system** supporting a single user.
 - The user can jump from one application to another, and each application uses the **same keyboard** for input and the **same screen** for output.





A Simple Example


- Each application needs to use the procedure echo,
- So it makes sense for it to be a **shared procedure** that is loaded into a portion of memory global to all applications.
- Thus, only a **single copy** of the echo procedure is used, **saving space**.





A Simple Example: Single Processor, Multiprogramming System

- P1 calls echo, but **interrupted** after first line (**chin=x**)
- P2 activated, calls echo and **runs till end** (**chin=chout=y**)
- P1 resumes, but the value is **overwritten** (prints y!!)
- **Problem:**
 - Shared global variables / procedure
- **Solution:**
 - Permit only one process to access the procedure at a time (works for multiprocessing systems as well)





A Simple Example: On a Multiprocessor

Process P1

-
- chin = getchar();
-
- chout = chin;
- putchar(chout);
-
-

Process P2

-
-
- chin = getchar();
- chout = chin;
-
- putchar(chout);
-





Enforce Single Access

- If we **enforce a rule** that only one process may enter the function at a time then:
 - P1 & P2 run on separate processors
 - P1 enters echo first
 - P2 tries to enter but is blocked
 - P1 completes execution
 - P2 resumes and executes echo





Race Condition

- A **race condition** occurs when
 - “Multiple processes or threads read and write data items hence the final result depends on the order of execution of the processes”
 - The output depends on who finishes the race last.
- E.g. Consider that a global variable “a” is accessed by Processes P1 and P2
 - P1 writes a=1
 - P2 writes a=2
 - The loser wins!





Operating System Concerns

- What **design and management issues** are raised by the existence of concurrency?
- The OS must
 - **Keep track** of various processes
 - **Allocate** and **de-allocate** resources
 - **Protect** the data and resources against interference by other processes
 - Manage **concurrency** properly





Process Interaction Scenarios

- Can be divided in **three categories**:
 - Processes **unaware** of each other
 - Competition
 - Processes **indirectly aware** of each other
 - Co-operation by sharing
 - Processes **directly aware** of each other
 - Co-operation by communication





Process Interaction: Unaware

- Example: **Multiprogramming**
 - Each process has different goal
- OS needs to be concerned about the **competition for resources**
 - P1 and P2 may compete for Printer
- **Potential Control Problems:**
 - Mutual Exclusion
 - Deadlock
 - Starvation





Process Interaction: Unaware

- **Mutual Exclusion:**
 - If P1 and P2 both need printer
 - Mutual exclusion should be enforced
 - Only one process should be allowed to use printer at a time
- **Critical Resource:**
 - Printer can be called a critical resource for this case
- **Critical Section:**
 - The part of code where printer is being accessed can be called critical section





Process Interaction: Unaware

- **Deadlock:**

- Occurs due to enforcement of **mutual exclusion**
- Consider two processes (**P1,P2**) and two resources (**R1,R2**)
- Process P1 is allocated R1 and P2 is allocated R2
- P1 requires R2 to complete and P2 requires R1 to complete --- **Deadlock!**





Process Interaction: Unaware

- **Starvation:**

- Occurs due to enforcement of **mutual exclusion**
- Consider that three processes **P1, P2 and P3** need periodic access to **resource R**
- If P1 is allotted R --- P2, P3 get delayed
- When P1 exits critical section, either P2 or P3 will be allowed to enter critical section
 - Suppose P3 is granted access
 - After P3, if P1 gets access to R again, P2 might suffer from **starvation**






Process Interaction: Unaware

- Hence, the **general structure** for each process should be as follows:

```
void P1 {  
    while(true) {  
        //Preceding Code  
        enterCritical(Ra);  
        //critical section  
        exitcritical(Ra);  
        //following Code  
    }  
}
```





Process Interaction: Indirect Awareness

- Multiple processes may have access to **shared variables / files / databases**
- **Integrity** needs to be ensured
- **Data is held on resources** e.g. memory or other storage
 - Same control problems occur – Mutual Exclusion, Deadlock, Starvation
 - **Difference** here is we need only write operations to be mutually exclusive





Process Interaction: Indirect Awareness

- One more requirement: **Data Coherence**
- Consider the code given which has to make sure that **a=b**. (a and b are initialized with same value)

P1 :

$a = a + 1;$

$b = b + 1;$

P2 :

$b = 2 * b;$

$a = 2 * a;$





Process Interaction: Indirect Awareness

- Assume that mutual exclusion is done on a and b (If P1 is updating a, P2 can not update a)
- If the processes execute as given below, then $a=b$ will not hold

```
a = a + 1;  
b = 2 * b;  
b = b + 1;  
a = 2 * a;
```


- So we need to put entire execution sequence in **critical section**





Process Interaction: Direct Awareness


- Processes co-operate by **communication**
- When processes cooperate by communication, they participate in a common effort that links all of the processes.
- The communication provides a way to synchronize, or coordinate, the various activities.





Process Interaction: Direct Awareness

- Because **nothing is shared** between processes in the act of passing messages, **mutual exclusion** is not a control requirement for this sort of cooperation.
- However, the problems of **deadlock** (P1 & P2 waiting for messages from each other) and **starvation** (P1 and P2 keep communicating, starving P3) are still present.





Process Interaction

Table 5.2 Process Interaction

Degree of Awareness	Relationship	Influence That One Process Has on the Other	Potential Control Problems
Processes unaware of each other	Competition	<ul style="list-style-type: none">• Results of one process independent of the action of others• Timing of process may be affected	<ul style="list-style-type: none">• Mutual exclusion• Deadlock (renewable resource)• Starvation
Processes indirectly aware of each other (e.g., shared object)	Cooperation by sharing	<ul style="list-style-type: none">• Results of one process may depend on information obtained from others• Timing of process may be affected	<ul style="list-style-type: none">• Mutual exclusion• Deadlock (renewable resource)• Starvation• Data coherence
Processes directly aware of each other (have communication primitives available to them)	Cooperation by communication	<ul style="list-style-type: none">• Results of one process may depend on information obtained from others• Timing of process may be affected	<ul style="list-style-type: none">• Deadlock (consumable resource)• Starvation

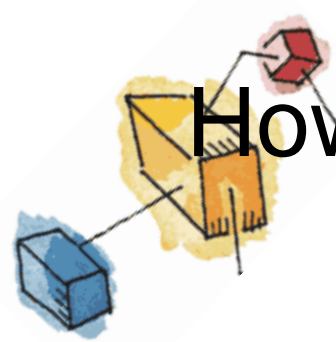




Requirements for Mutual Exclusion

- Only **one process at a time** is allowed in the **critical section** for a resource
- A process that halts in its **noncritical section** must do so without interfering with other processes
- No **deadlock** or **starvation**
- A process must not be **delayed** access to a critical section when there is no other process using it
- **No assumptions** are made about relative process speeds or number of processes
- A process remains inside its critical section for a **finite time** only

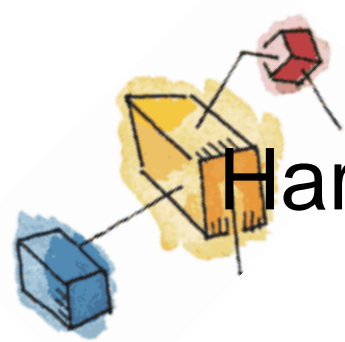




How to achieve the Requirements of Mutual Exclusion?

- Three ways possible
 1. Leave the responsibility to the processes
 - High overhead
 2. Use special purpose machine instructions (hardware support)
 - Reduces overhead
 3. Support of OS
 - Most popular





Hardware Support for Mutual Exclusion

- Two methods:

1. Interrupt Disabling

2. Special Machine Instructions





Disabling Interrupts

- **Uniprocessors** only allow interleaving
- A process runs until it invokes an operating system service or until it is interrupted
- **Disabling interrupts** guarantees mutual exclusion
- **Issues:**
 - Execution efficiency is degraded as interleaving gets limited.
 - Will not work in multiprocessor architecture

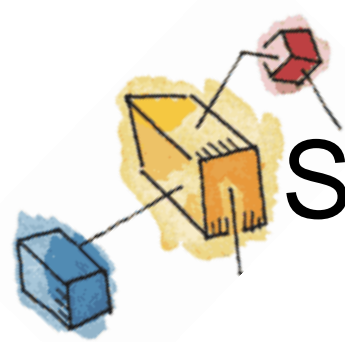




Pseudo-Code

```
while (true) {  
    /* disable interrupts */;  
    /* critical section */;  
    /* enable interrupts */;  
    /* remainder */;  
}
```

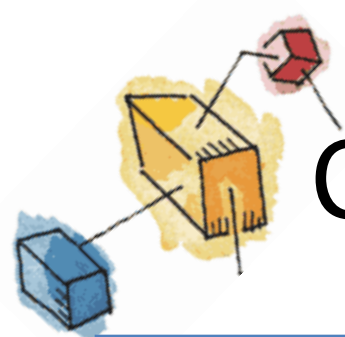




Special Machine Instructions

- Compare&Swap Instruction
 - also called a “compare and exchange instruction”
- Exchange Instruction

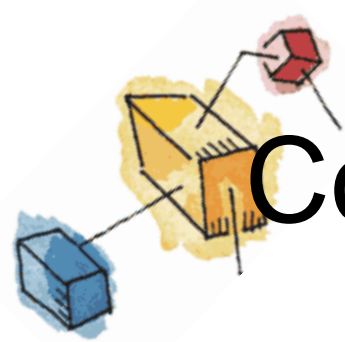




Compare&Swap Instruction

```
int compare_and_swap (int *word,  
    int testval, int newval)  
{  
    int oldval;  
    oldval = *word;  
    if (oldval == testval) *word = newval;  
    return oldval;  
}
```





Compare&Swap Instruction

- This version of the instruction checks a **memory location** (*word) against a **test value** (testval).
- If the memory location's current value is testval, it is replaced with newval;
 - otherwise it is left unchanged.
- The **old memory value** is always returned;
 - thus, the memory location has been updated if the returned value is the same as the test value.
- This is an **atomic instruction**





Mutual Exclusion

```
/* program mutualexclusion */
const int n = /* number of processes */;
int bolt;
void P(int i)
{
    while (true) {
        while (compare_and_swap(bolt, 0, 1) == 1)
            /* do nothing */;
        /* critical section */;
        bolt = 0;
        /* remainder */;
    }
}
void main()
{
    bolt = 0;
    parbegin (P(1), P(2), ..., P(n));
}
```

(a) Compare and swap instruction

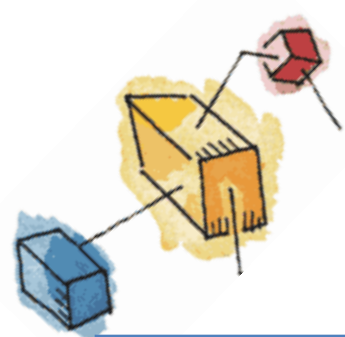




Mutual Exclusion

- Bolt is initialized to 0
- First call (P1) would be $C\&S(0,0,1)$
 - Word and testval same so word gets updated
 - Return value is 0, hence P1 can enter critical section (while condition false)
- P2 calls $C\&S(1,0,1)$
 - No change
 - Return value is 1, hence P2 waits (while condition true)





Exchange instruction

```
void exchange (int register, int
memory)
{
    int temp;
    temp = memory;
    memory = register;
    register = temp;
}
```

- Exchange the contents of register with the content of memory location





Mutual Exclusion

```
/* program mutualexclusion */
int const n = /* number of processes**/;
int bolt;
void P(int i)
{
    int keyi = 1;
    while (true) {
        do exchange (keyi, bolt)
        while (keyi != 0);
        /* critical section */;
        bolt = 0;
        /* remainder */;
    }
}
void main()
{
    bolt = 0;
    parbegin (P(1), P(2), ..., P(n));
}
```

(b) Exchange instruction

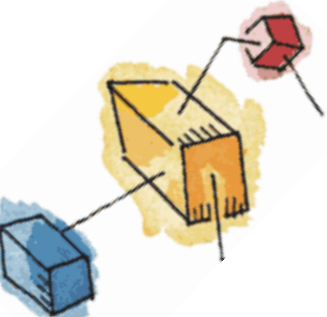




Mutual Exclusion

- For process P1,
 - $key_i=1$, $exchange(1,0)$ is called
 - $key_i=0$, $bolt=1$
 - P1 enters critical section
- For process P2,
 - $key_i=1$, $exchange(1,1)$ is called
 - No change, P2 waits





Hardware Mutual Exclusion: Advantages

- **Applicable to** any number of processes on either a single processor or multiple processors sharing main memory
- It is **simple** and therefore easy to verify
- It can be used to support **multiple critical sections**





Hardware Mutual Exclusion: Disadvantages

- **Busy-waiting** consumes processor time
- **Starvation is possible** when a process leaves a critical section and more than one process is waiting.
 - Some process could indefinitely be denied access.
- **Deadlock is possible**
 - P1 runs and calls C&S
 - Interrupted due to high priority process P2
 - P2 runs but needs the resource occupied by P1





Semaphore

- OS and Programming language based mechanism
- Two or more processes co-operate using **signals**
 - A process has to stop at a place until it receives a signal
 - Special variables – **Semaphores** are used for this purpose
- **Operations on Semaphores** (all of which are atomic):
 - initialize,
 - Decrement (`semWait`)
 - increment. (`semSignal`)





Operations on Semaphores

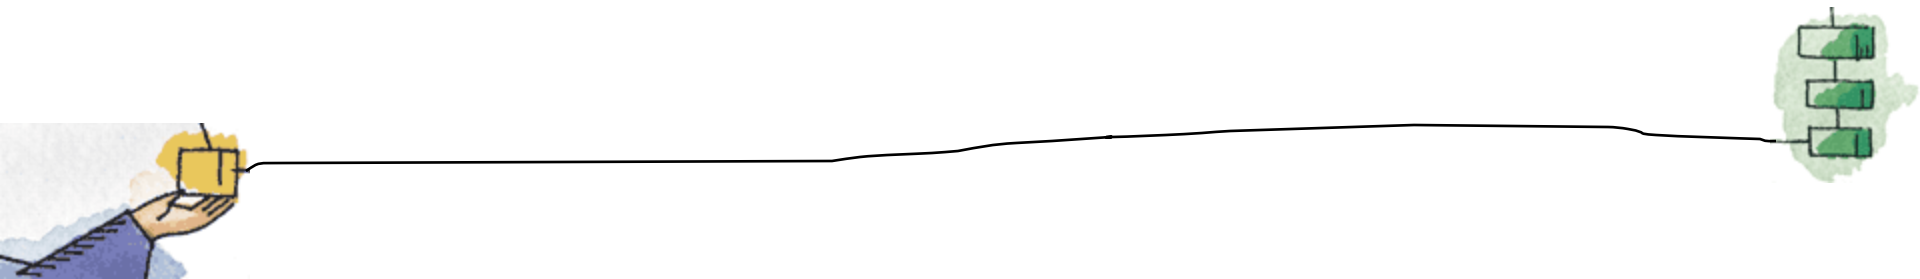
- A semaphore may be **initialized** to a **nonnegative integer** value.
- The **semWait** operation **decrements** the semaphore value.
 - If the value becomes **negative**, then the process executing the semWait is **blocked**.
 - **Otherwise**, the process **continues** execution.





Operations on Semaphores

- The **semSignal** operation **increments** the semaphore value.
 - If the resulting value is **less than or equal to zero**, then a process blocked by a semWait operation, if any, is **unblocked**.
- This type of semaphores are called **Counting / General semaphores**





Semaphore Primitives

```
struct semaphore {
    int count;
    queueType queue;
};
void semWait(semaphore s)
{
    s.count--;
    if (s.count < 0) {
        /* place this process in s.queue */;
        /* block this process */;
    }
}
void semSignal(semaphore s)
{
    s.count++;
    if (s.count <= 0) {
        /* remove a process P from s.queue */;
        /* place process P on ready list */;
    }
}
```



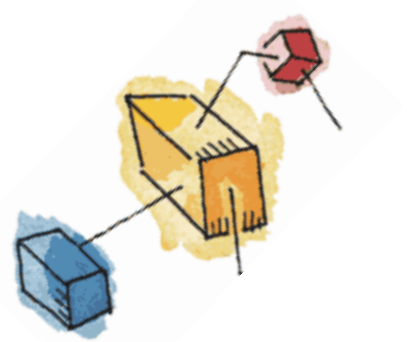


Figure 5.3 A Definition of Semaphore Primitives



Binary Semaphores

- **Restricted version** of counting semaphores
- **Operations on Semaphores** (all of which are atomic):
 - Initialize
 - Decrement (`semWaitB`)
 - increment. (`semSignalB`)

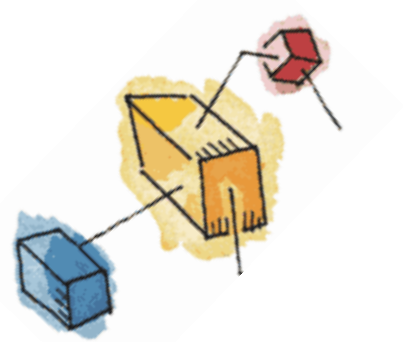




Binary Semaphores

- A binary semaphore may be **initialized** to either **zero** or **one**.
- The **semWaitB** operation **checks** the semaphore value
 - If the value is **zero**, process is **blocked**
 - **Otherwise**, value is **changed to zero** and the process **continues** execution.

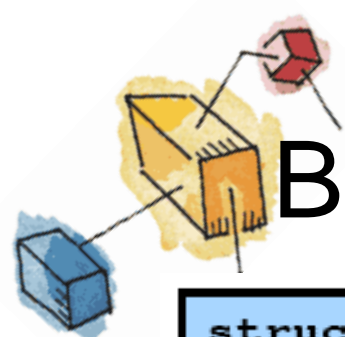




Binary Semaphores

- The **semSignalB** operation **checks** if any processes are **blocked** (semaphore value zero)
 - If **blocked**, one process is **unblocked**
 - If **no processes** are **blocked**, value is **set to One**.





Binary Semaphore Primitives

```
struct binary_semaphore {
    enum {zero, one} value;
    queueType queue;
};

void semWaitB(binary_semaphore s)
{
    if (s.value == one)
        s.value = zero;
    else {
        /* place this process in s.queue */;
        /* block this process */;
    }
}

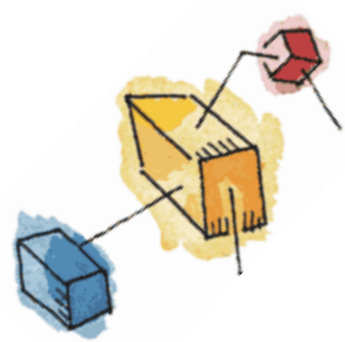
void semSignalB(semaphore s)
{
    if (s.queue is empty())
        s.value = one;
    else {
        /* remove a process P from s.queue */;
        /* place process P on ready list */;
    }
}
```

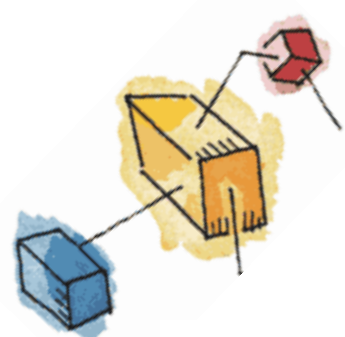


Figure 5.4 A Definition of Binary Semaphore Primitives

Strong/Weak Semaphore

- A **queue** is used to hold processes waiting on the semaphore
 - In what order are processes removed from the queue?
- **Strong Semaphores** use **FIFO**
- **Weak Semaphores** don't specify the order of removal from the queue



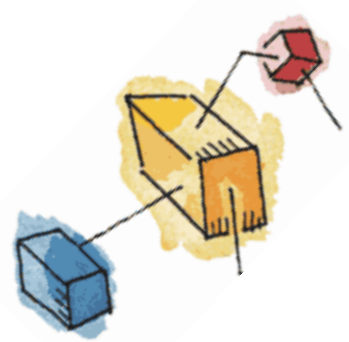


Mutual Exclusion Using Semaphores

```
/* program mutualexclusion */
const int n = /* number of processes */;
semaphore s = 1;
void P(int i)
{
    while (true) {
        semWait(s);
        /* critical section */;
        semSignal(s);
        /* remainder */;
    }
}
void main()
{
    parbegin (P(1), P(2), . . . , P(n));
}
```

Figure 5.6 Mutual Exclusion Using Semaphores





Mutual Exclusion Using Semaphores

- All processes need the same resource, hence each process contains a critical section to access the resource
- First process executing semWait will be allowed to enter critical section, rest of the processes will be blocked
- When a process exits critical section, semaphore value is incremented, so any one blocked process will be allowed to enter critical section
- Can the same solution work for multiple critical sections?





Mutual Exclusion Using Semaphores

- If the value of **s is positive**,
 - Equal to number of processes that can issue semWait and access critical section
- If the value of **s is zero**,
 - Next process that issues semWait will be blocked
- If the value of **s is negative**,
 - Equal to the number of processes waiting





Producer/Consumer Problem

- **General Situation:**

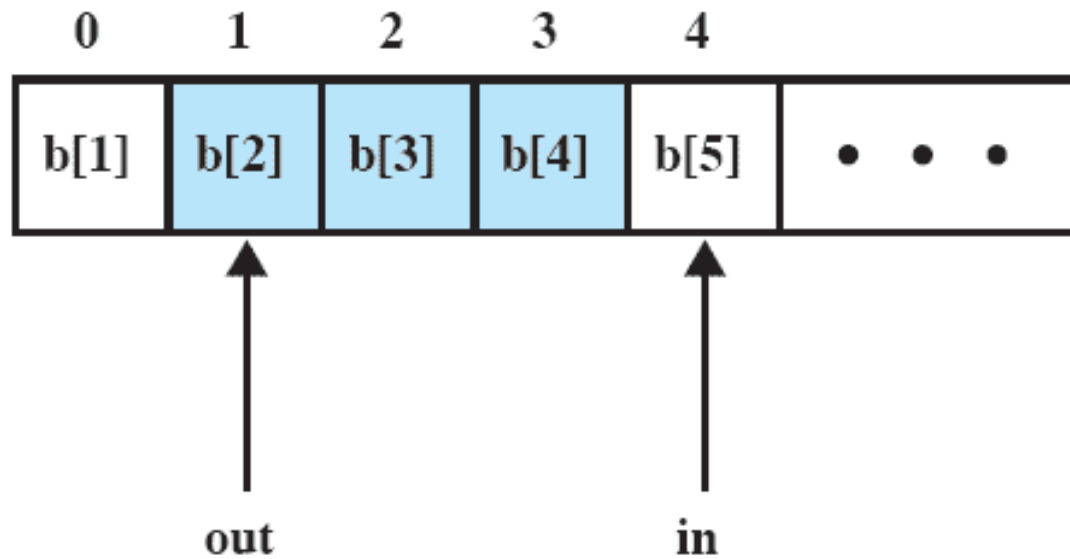
- One or more **producers** are generating data and placing these in a buffer
- A **single consumer** is taking items out of the buffer one at time
- **Only one** producer or consumer may access the buffer at any one time

- **The Problem:**

- Ensure that the Producer can't add data into **full buffer** and consumer can't remove data from **empty buffer**



Buffer



Note: shaded area indicates portion of buffer that is occupied


Figure 5.8 Infinite Buffer for the Producer/Consumer Problem



Functions

- Assume an infinite buffer ***b*** with a linear array of elements

Producer	Consumer
<pre>while (true) { /* produce item v */ b[in] = v; in++; }</pre>	<pre>while (true) { while (in <= out) /*do nothing */; w = b[out]; out++; /* consume item w */ }</pre>



Incorrect Solution

```
/* program producerconsumer */
int n;
binary_semaphore s = 1, delay = 0;
void producer()
{
    while (true) {
        produce();
        semWaitB(s);
        append();
        n++;
        if (n==1) semSignalB(delay);
        semSignalB(s);
    }
}
void consumer()
{
    semWaitB(delay);
    while (true) {
        semWaitB(s);
        take();
        n--;
        semSignalB(s);
        consume();
        if (n==0) semWaitB(delay);
    }
}
void main()
{
    n = 0;
    parbegin (producer, consumer);
}
```

Possible Scenario

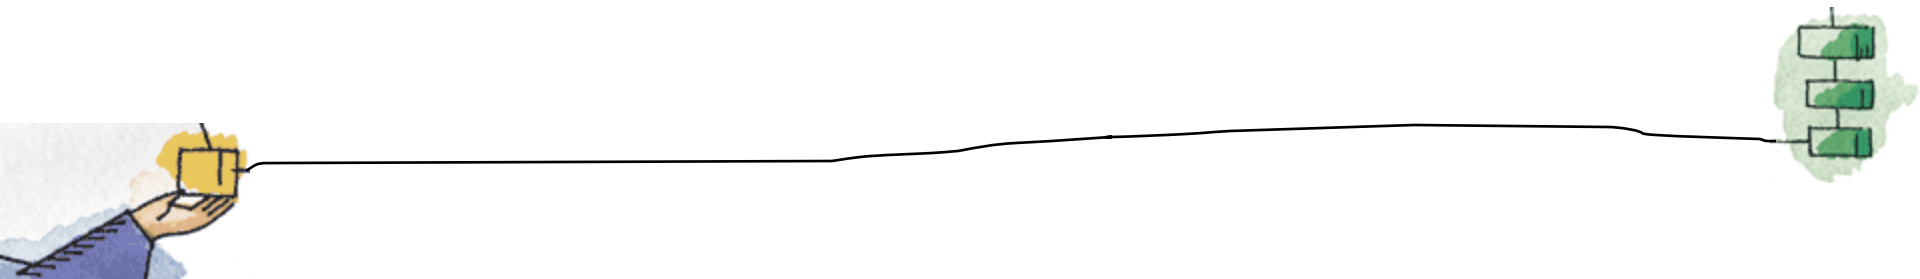
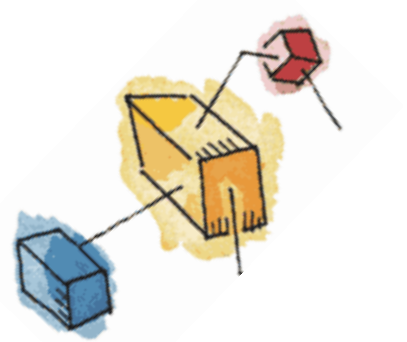
Table 5.4 Possible Scenario for the Program of Figure 5.9

	Producer	Consumer	s	n	Delay
1			1	0	0
2	semWaitB(s)		0	0	0
3	n++		0	1	0
4	if (n==1) (semSignalB(delay))		0	1	1
5	semSignalB(s)		1	1	1
6		semWaitB(delay)	1	1	0
7		semWaitB(s)	0	1	0
8		n--	0	0	0
9		semSignalB(s)	1	0	0
10	semWaitB(s)		0	0	0
11	n++		0	1	0
12	if (n==1) (semSignalB(delay))		0	1	1
13	semSignalB(s)		1	1	1
14		if (n==0) (semWaitB(delay))	1	1	1
15		semWaitB(s)	0	1	1
16		n--	0	0	1
17		semSignalB(s)	1	0	1
18		if (n==0) (semWaitB(delay))	1	0	0
19		semWaitB(s)	0	0	0
20		n--	0	-1	0
21		semiSignlaB(s)	1	-1	0

NOTE: White areas represent the critical section controlled by semaphore s.

Solution?

- What will happen if the conditional statement is moved inside critical section?
 - Deadlock!
 - Consumer would block in critical section
 - Producer can't enter critical section to unblock consumer



Correct Solution

```
/* program producerconsumer */
int n;
binary_semaphore s = 1, delay = 0;
void producer()
{
    while (true) {
        produce();
        semWaitB(s);
        append();
        n++;
        if (n==1) semSignalB(delay);
        semSignalB(s);
    }
}
void consumer()
{
    int m; /* a local variable */
    semWaitB(delay);
    while (true) {
        semWaitB(s);
        take();
        n--;
        m = n;
        semSignalB(s);
        consume();
        if (m==0) semWaitB(delay);
    }
}
void main()
{
    n = 0;
    parbegin (producer, consumer);
}
```




Using Counting Semaphores

```
/* program producerconsumer */
semaphore n = 0, s = 1;
void producer()
{
    while (true) {
        produce();
        semWait(s);
        append();
        semSignal(s);
        semSignal(n);
    }
}
void consumer()
{
    while (true) {
        semWait(n);
        semWait(s);
        take();
        semSignal(s);
        consume();
    }
}
void main()
{
    parbegin (producer, consumer);
}
```

Figure 5.11 A Solution to the Infinite-Buffer Producer/Consumer Problem Using Semaphores



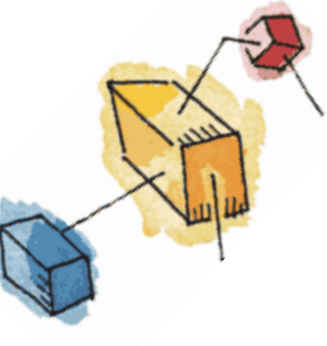


Cases to consider

- To check the working of solution, we can test the following scenarios:
 - Producer runs before consumer
 - Consumer runs before producer
 - Multiple producers run and then consumer runs
- Will the interchange of first two lines in consumer code create any issue?



Bounded Buffer



Block on:	Unblock on:
Producer: insert in full buffer	Consumer: item inserted
Consumer: remove from empty buffer	Producer: item removed







Figure 5.12 Finite Circular Buffer for the Producer/Consumer Problem





Bounded Buffer


- Three semaphores needed:
 - S: Mutual Exclusion
 - E: Keep track of empty spaces
 - N: Number of items
- What should be the initial values for semaphores?



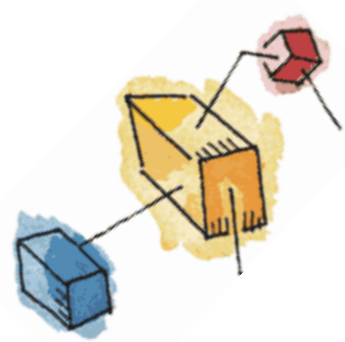


Bounded Buffer

```
/* program boundedbuffer */
const int sizeofbuffer = /* buffer size */;
semaphore s = 1, n = 0, e = sizeofbuffer;
void producer()
{
    while (true) {
        produce();
        semWait(e);
        semWait(s);
        append();
        semSignal(s);
        semSignal(n);
    }
}
void consumer()
{
    while (true) {
        semWait(n);
        semWait(s);
        take();
        semSignal(s);
        semSignal(e);
        consume();
    }
}
void main()
{
    parbegin (producer, consumer);
}
```



Functions in a Bounded Buffer

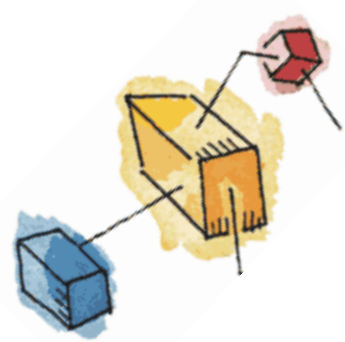


Producer	Consumer
<pre>while (true) { /* produce item v */ while ((in + 1) % n == out) /* do nothing */; b[in] = v; in = (in + 1) % n }</pre>	<pre>while (true) { while (in == out) /* do nothing */; w = b[out]; out = (out + 1) % n; /* consume item w */ }</pre>



Monitors

- **Semaphores** are **difficult to control** as semWait and SemSignal calls can be scattered throughout the program.
- The **Monitor** is a **programming-language construct** that provides equivalent functionality to that of semaphores and that is easier to control.
- **Monitor contains:**
 - One or more procedures
 - Initialization sequence
 - Local data





Chief characteristics

1. **Local data** variables are accessible by monitor's procedures
 2. A **process enters monitor** by invoking one of its procedures
 3. Only **one process** may be executing in the monitor **at a time**
 - Any other process invoking monitor is blocked
- 1st and 2nd are **like OOP**
 - 3rd provides **mutual exclusion**



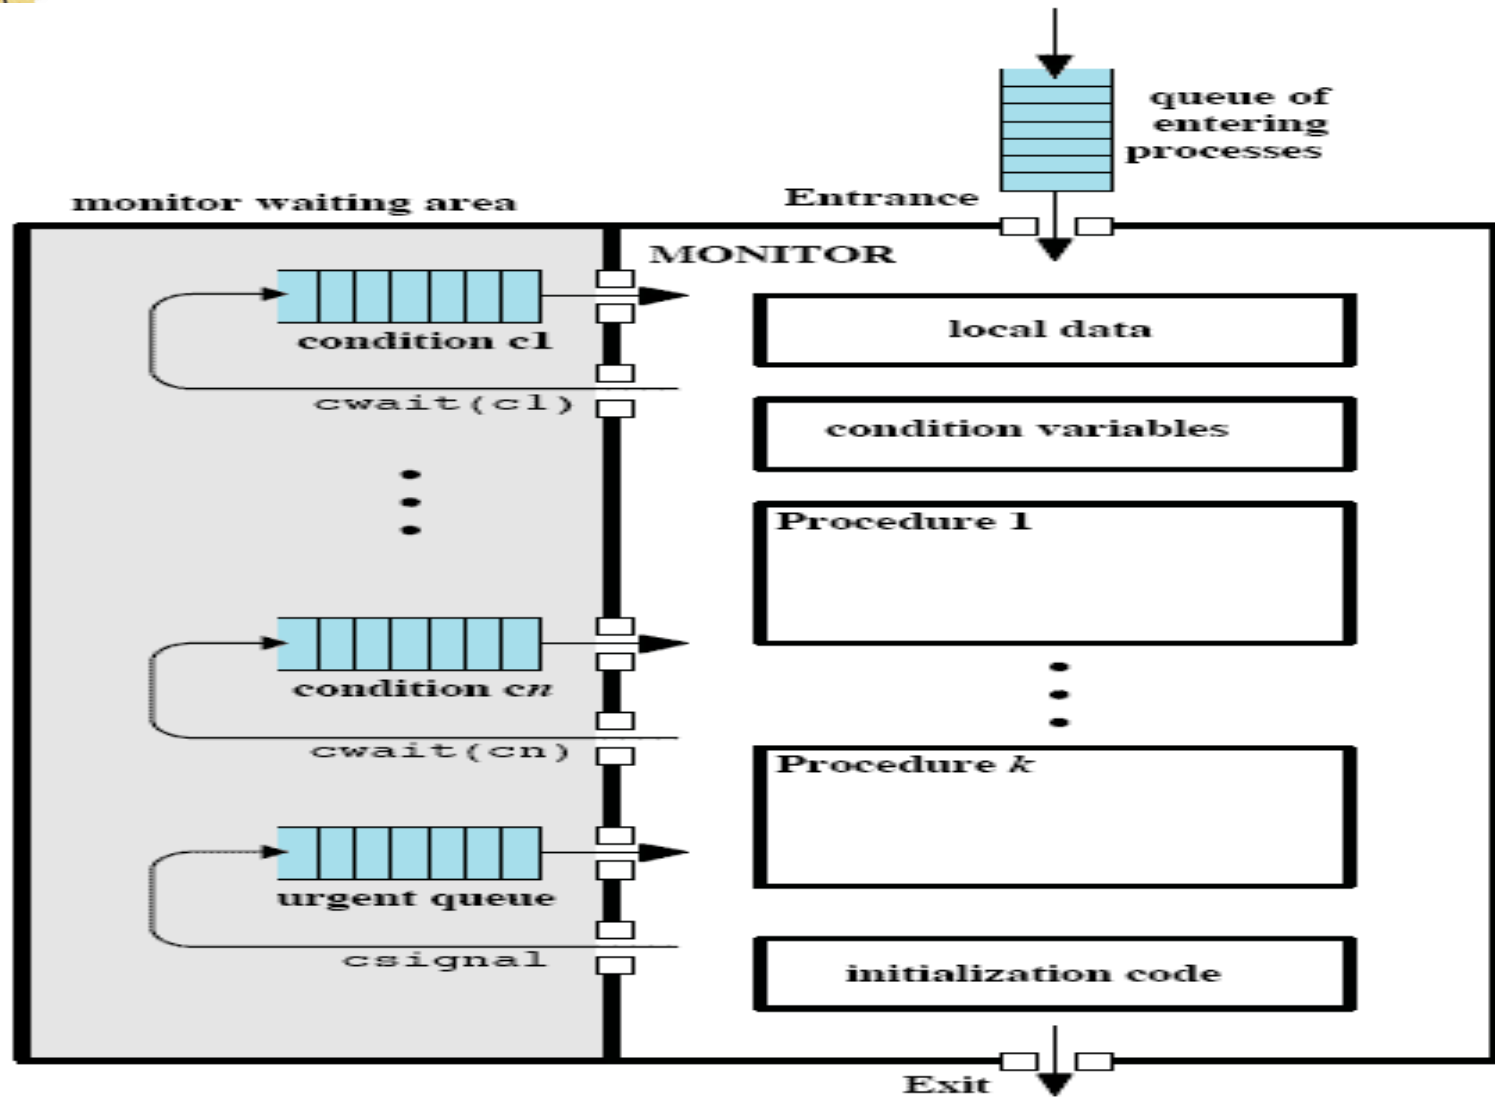


Synchronization

- Synchronisation achieved by **condition variables** within a monitor
- **Monitor Functions:**
 - **Cwait(c):** Suspend execution of the calling process on condition *c*
 - **Csignal(c):** Resume execution of some process blocked after a cwait on the same condition
- **Difference with Semaphores:**
 - If a process in monitor signals and no task is waiting on condition variable, the signal is lost



Structure of a Monitor





Structure of a Monitor

- Although a process can enter the monitor by invoking any of its procedures, we can think of the monitor as having a single entry point that is guarded so that only one process may be in the monitor at a time.
 - Other processes that attempt to enter the monitor join a queue of processes blocked waiting for monitor availability.
- Once a process is in the monitor, it may temporarily block itself on condition x by issuing `cwait(x)`





Structure of a Monitor

- It is then placed in a **queue of processes waiting to re-enter the monitor** when the condition changes, and resume execution at the point in its program following the `cwait(x)` call.
- If a **process that is executing in the monitor** detects a change in the condition variable `x`, it issues **`csignal(x)`**, which alerts the corresponding condition queue that the condition has changed.





Bounded Buffer Solution Using Monitor

```
/* program producerconsumer */
monitor boundedbuffer;
char buffer [N];                                /* space for N items */
int nextin, nextout;                             /* buffer pointers */
int count;                                       /* number of items in buffer */
cond notfull, notempty;                        /* condition variables for synchronization */

void append (char x)
{
    if (count == N) cwait(notfull);             /* buffer is full; avoid overflow */
    buffer[nextin] = x;
    nextin = (nextin + 1) % N;
    count++;
    /* one more item in buffer */
    csignal(notempty);                          /* resume any waiting consumer */
}

void take (char x)
{
    if (count == 0) cwait(notempty);             /* buffer is empty; avoid underflow */
    x = buffer[nextout];
    nextout = (nextout + 1) % N;
    count--;
    /* one fewer item in buffer */
    csignal(notfull);                          /* resume any waiting producer */
}

/* monitor body */
{
    nextin = 0; nextout = 0; count = 0;          /* buffer initially empty */
}
```



Bounded Buffer Solution Using Monitor


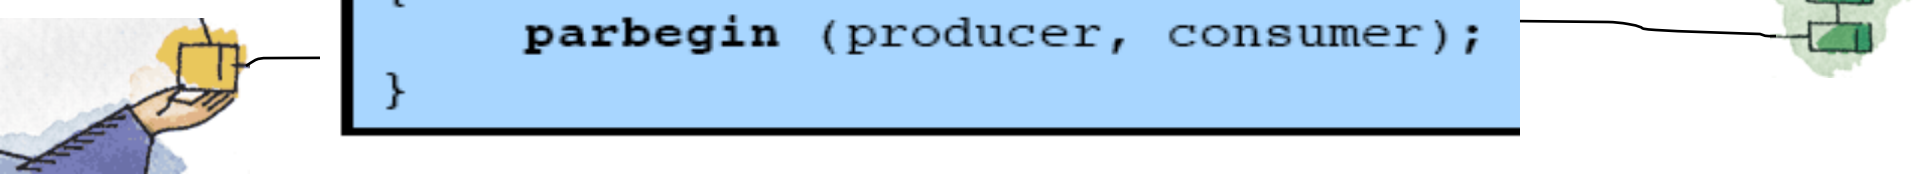
- **Notfull is true** when there is a room to add at least once character to the buffer
- **Notempty is true** when there is at least one character in the buffer
- **Producer** can add characters to the buffer by means of **append** procedure only
- **Consumer** can remove characters from the buffer by means of **take** procedure only
- **“Monitor itself enforces mutual exclusion, programmer has to manage synchronization only”**
- **“Semaphores require programmer to manage both”**





Solution Using Monitor

```
void producer()  
{  
    char x;  
    while (true) {  
        produce(x);  
        append(x);  
    }  
}  
void consumer()  
{  
    char x;  
    while (true) {  
        take(x);  
        consume(x);  
    }  
}  
void main()  
{  
    parbegin (producer, consumer);  
}
```





Message Passing

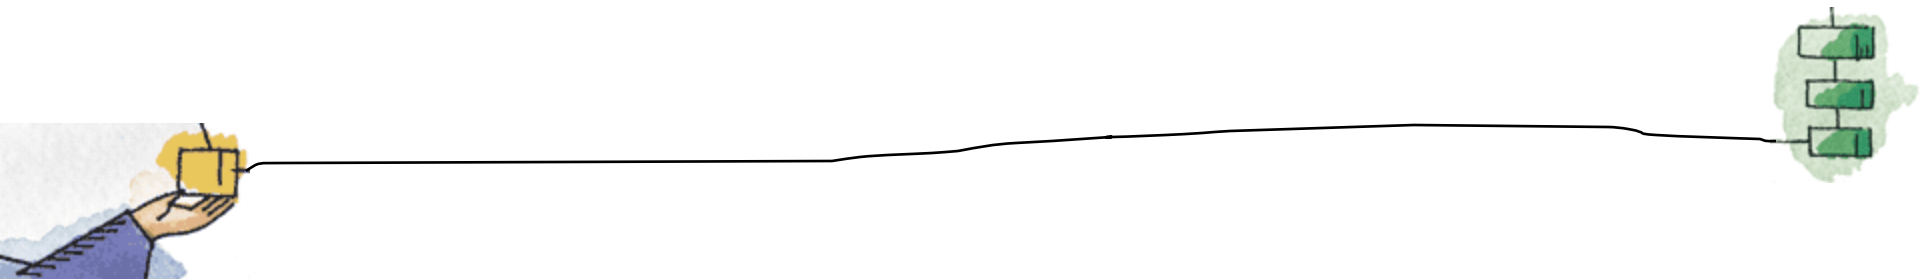
- When **processes interact** with one another, two fundamental requirements must be satisfied:
 - synchronization and
 - communication.
- The actual function of message passing is normally provided in the form of a pair of primitives:
 - **send (destination, message)**
 - **receive (source, message)**

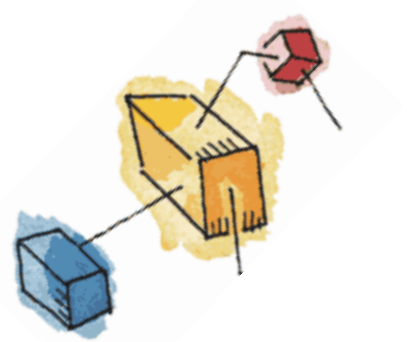




Synchronization


- “Communication requires synchronization”
 - Sender must send before receiver can receive
- What happens to a process after it issues a send or receive primitive?
 - Sender and receiver may or may not be blocking (waiting for message)






Message Passing

- Using these functions, **three useful combinations** are commonly used:
 1. **Blocking** Send, **Blocking** Receive
 2. **Non Blocking** Send, **Blocking** Receive
 3. **Non Blocking** Send, **Non Blocking** Receive



Blocking send, Blocking receive

- Both sender and receiver are blocked until message is delivered
- Known as a *rendezvous*
- Allows for *tight synchronization* between processes.






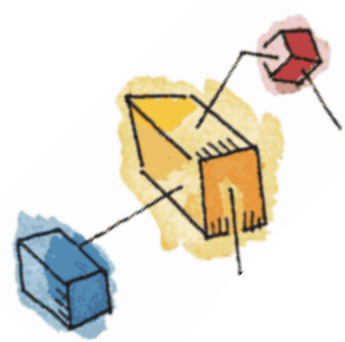
Non-blocking Send


- More natural for many concurrent programming tasks
- Nonblocking send, blocking receive
 - Sender continues on
 - Receiver is blocked until the requested message arrives
- Nonblocking send, nonblocking receive
 - Neither party is required to wait



Addressing


- Sending process need to be able to specify which process should receive the message
- Two Types:
 - Direct addressing
 - Indirect Addressing






Direct Addressing

- **Send primitive** includes a **specific identifier** of the destination process
- **Receive primitive** could know ahead of time which process a message is expected OR
- **Receive primitive** could use source parameter to return a value when the receive operation has been performed





Indirect addressing

- Messages are sent to a **shared data structure** consisting of queues
- Queues are called *mailboxes*
- One process sends a message to the mailbox and the other process picks up the message from the mailbox
- **Advantage:**
 - Decoupling of sender and receiver





Indirect addressing

- **Four Types:**

1. One to One -- e.g. private communication
2. One to Many – e.g. broadcast
3. Many to One – e.g. client-server
4. Many to Many – e.g. multiple clients, multiple servers



Indirect Process Communication

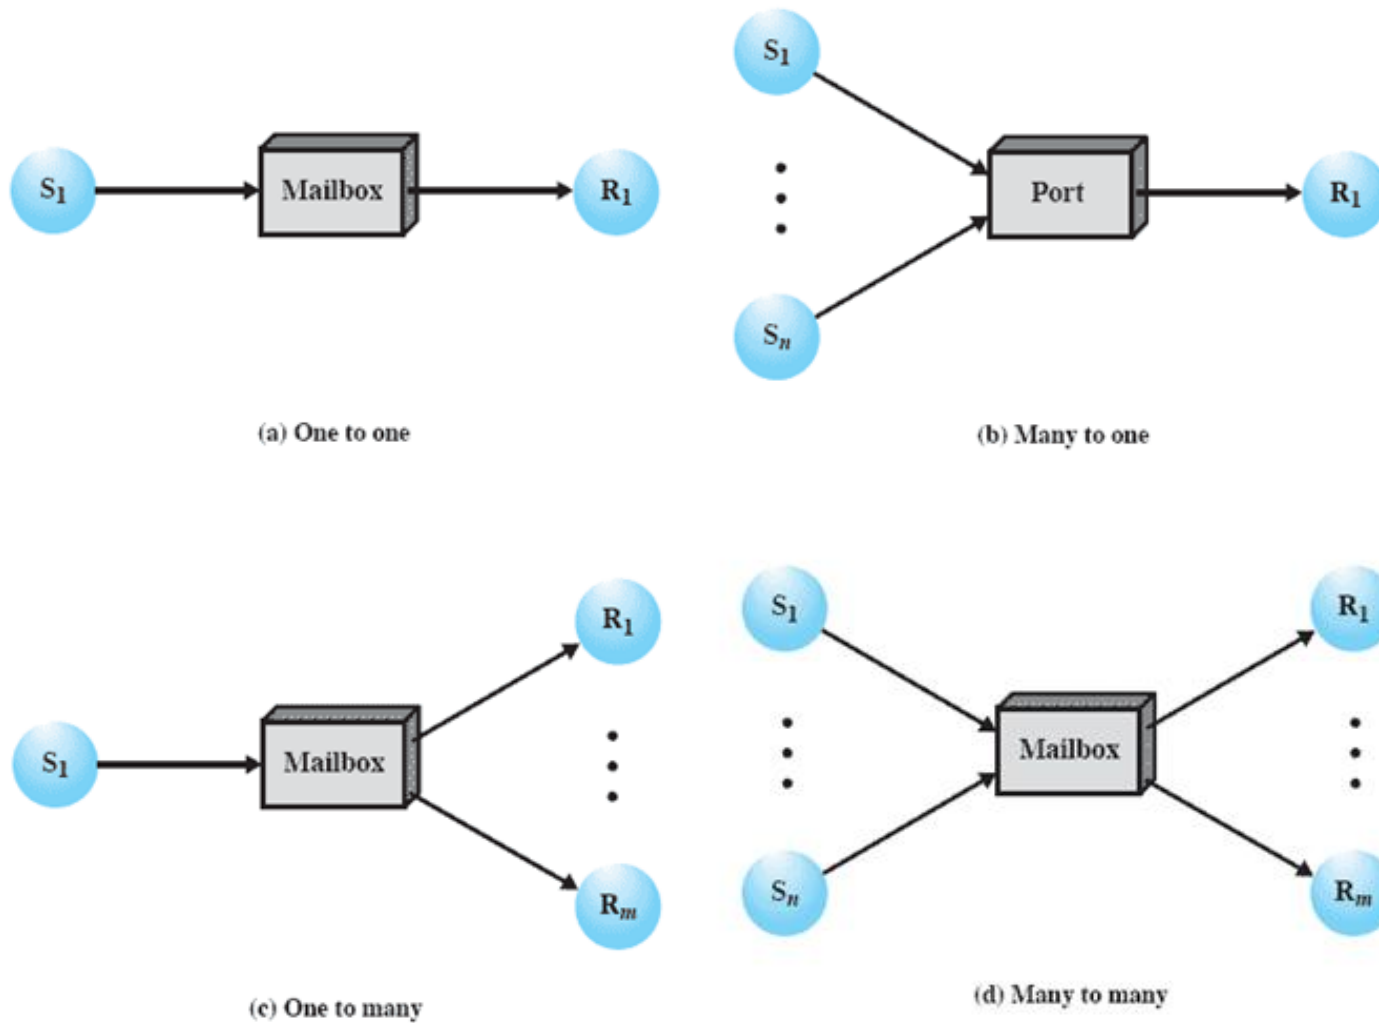


Figure 5.18 Indirect Process Communication



General Message Format




Figure 5.19 General Message Format





Mutual Exclusion Using Messages

- The approach can be designed as follows:
 - Attempt to receive the message
 - If MailBox is empty
 - block the process
 - Else
 - Enter Critical Section
 - Extract Message
 - Place Message in box for other processes
- Nonblocking send, blocking receive approach






Mutual Exclusion Using Messages

```
/* program mutualexclusion */
const int n = /* number of processes */;
void P(int i)
{
    message msg;
    while (true) {
        receive (box, msg);
        /* critical section    */;
        send (box, msg);
        /* remainder    */;
    }
}
void main()
{
    create mailbox (box);
    send (box, null);
    parbegin (P(1), P(2), . . ., P(n));
}
```

Figure 5.20 Mutual Exclusion Using Messages





Mutual Exclusion Using Messages

- We assume the use of the **blocking receive** primitive and the **non-blocking send** primitive.
- This assumes that if more than one process performs the receive operation concurrently, then
 - If there is a message, it is delivered to only one process and the others are blocked, or
 - If the message queue is empty, all processes are blocked; when a message is available, only one blocked process is activated and given the message.






Producer/Consumer Messages

```
const int
    capacity = /* buffering capacity */ ;
    null = /* empty message */ ;
int i;
void producer()
{   message pmsg;
    while (true) {
        receive (mayproduce, pmsg);
        pmsg = produce();
        send (mayconsume, pmsg);
    }
}
void consumer()
{   message cmsg;
    while (true) {
        receive (mayconsume, cmsg);
        consume (cmsg);
        send (mayproduce, null);
    }
}

void main()
{
    create_mailbox (mayproduce);
    create_mailbox (mayconsume);
    for (int i = 1; i <= capacity; i++) send (mayproduce, null);
    parbegin (producer, consumer);
}
```





Producer/Consumer Messages

- Two mailboxes are used.
 - As the **producer** generates data, it is sent as messages to the mailbox **mayconsume**.
 - As long as there is at least one message in that mailbox, the **consumer** can consume.
- Hence **mayconsume serves as the buffer**; the data in the buffer are organized as a queue of messages.






Producer/Consumer Messages

- The “size” of the buffer is determined by the global variable capacity.
- Initially, the mailbox `mayproduce` is filled with a number of null messages equal to the capacity of the buffer.
- The `number of messages` in `mayproduce` shrinks with each production and grows with each consumption.





Readers/Writers Problem



- A data area is shared among many processes
 - Some processes only read the data area, some only write to the area
- Conditions to satisfy:
 1. Multiple readers may read the file at once.
 2. Only one writer at a time may write
 3. If a writer is writing to the file, no reader may read it.





Readers have Priority



```
/* program readersandwriters */
int readcount;
semaphore x = 1, wsem = 1;
void reader()
{
    while (true) {
        semWait (x);
        readcount++;
        if (readcount == 1) semWait (wsem);
        semSignal (x);
        READUNIT();
        semWait (x);
        readcount--;
        if (readcount == 0) semSignal (wsem);
        semSignal (x);
    }
}
void writer()
{
    while (true) {
        semWait (wsem);
        WRITEUNIT();
        semSignal (wsem);
    }
}
void main()
{
    readcount = 0;
    parbegin (reader, writer);
}
```





Writers have Priority

```
/* program readersandwriters */
int  readcount, writecount;
semaphore x = 1, y = 1, z = 1, wsem = 1, rsem = 1;
void reader()
{
    while (true) {
        semWait (z);
        semWait (rsem);
        semWait (x);
        readcount++;
        if (readcount == 1) semWait (wsem);
        semSignal (x);
        semSignal (rsem);
        semSignal (z);
        READUNIT();
        semWait (x);
        readcount--;
        if (readcount == 0) semSignal (wsem);
        semSignal (x);
    }
}
```





Writers have Priority

```
void writer ()
{
    while (true) {
        semWait (y);
        writecount++;
        if (writecount == 1) semWait (rsem);
        semSignal (y);
        semWait (wsem);
        WRITEUNIT();
        semSignal (wsem);
        semWait (y);
        writecount--;
        if (writecount == 0) semSignal (rsem);
        semSignal (y);
    }
}

void main()
{
    readcount = writecount = 0;
    parbegin (reader, writer);
}
```

